


```
In [ ]: # import pandas as pd, numpy as np
pd.set_option('display.width',150)
pd.set_option('display.max_columns',20)
df = pd.read_csv("GCO_Sample_Data_2022.csv")
df.columns
```

```
In [5]: df.plot(x='Big4',y='Size')
```

```
Out[5]: <Axes: xlabel='Big4'>
```



```
In [6]: col1, col2 = "Size", "Big4"
corr = df[col1].corr(df[col2])
print ("Correlation between ", col1, " and ", col2, "is: ", round(corr, 2)) #
Correlation between Size and Big4 is: 0.41
```

```
In [7]: df.groupby(['GCO']).agg({'Big4':'mean','Loss':lambda x: list(x)})
```

```
Out[7]:
```

	Big4	Loss
GCO		
0	0.780877	[0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1, 1, ...]
1	0.566038	[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, ...]

```
In [8]: df.isnull().sum() # Return null values
df['GCO'].value_counts()
```

```
Out[8]: 0    4947
        1     53
        Name: GC0, dtype: int64
```

```
In [9]: # Replace according to the mapping table provided above
df['goingconcern'] = df['GC0'].replace({1 : 'GC_Warning', 0: 'Opinion_Clean'})
df.goingconcern.value_counts()
```

```
Out[9]: Opinion_Clean    4947
        GC_Warning       53
        Name: goingconcern, dtype: int64
```

```
In [10]: df['GC0'].describe()
```

```
Out[10]: count    5000.000000
         mean      0.010600
         std       0.102419
         min       0.000000
         25%       0.000000
         50%       0.000000
         75%       0.000000
         max       1.000000
         Name: GC0, dtype: float64
```

```
In [11]: # Approach 1 - Using Pandas
print(f"Using Pandas:")
print(f"mean: {df.GC0.mean()}")
print(f"standard_deviation: {df.GC0.std()}")
print(f"minimum: {df.GC0.min()}")
print(f"25th_percentile: {df.GC0.quantile(0.25)}")
print(f"50th_percentile: {df.GC0.quantile(0.50)}")
print(f"75th_percentile: {df.GC0.quantile(0.75)}")
print(f"maximum: {df.GC0.max()}\n")
```

```
Using Pandas:
mean: 0.0106
standard_deviation: 0.10241942173040403
minimum: 0
25th_percentile: 0.0
50th_percentile: 0.0
75th_percentile: 0.0
maximum: 1
```

```
In [12]: # Approach 2 - Using NumPy
print(f"Using NumPy:")
print(f"mean: {np.mean(df.GC0)}")
print(f"standard_deviation: {np.std(df.GC0, ddof = 1)}")
print(f"minimum: {np.min(df.GC0)}")
print(f"25th_percentile: {np.percentile(df.GC0, 25)}")
print(f"50th_percentile: {np.percentile(df.GC0, 50)}")
print(f"75th_percentile: {np.percentile(df.GC0, 75)}")
print(f"maximum: {np.max(df.GC0)}\n")
```


```
Using NumPy:
mean: 0.0106
standard_deviation: 0.10241942173040403
minimum: 0
25th_percentile: 0.0
50th_percentile: 0.0
75th_percentile: 0.0
maximum: 1
```

```
In [13]: lower_bound = np.mean(df['Size'][df.GC0 == 1])
         upper_bound = np.mean(df['Size'][df.GC0 == 0])
```


```
In [14]: print(f"lower: {lower_bound}")
         print(f"upper: {upper_bound}")
```

```
lower: 4.109821029433963
upper: 6.928418017709723
```

```
In [15]: import seaborn as sns, matplotlib.pyplot as plt
         sns.histplot(df.Size)
         plt.show()
```




```
In [16]: plt.figure(figsize = (15, 5))
         sns.boxplot(data = df, x = 'Loss', y = 'Size')
         plt.show()
```



```
In [19]: xvars = ['Tenure', 'FeeRatio', 'FirmAge', 'ZScore', 'Size', 'Loss', 'Leverage',
                'ChLeverage', 'OpCF', 'Finance', 'Invest', 'Returns', 'Beta', 'Volatility']
# train an XGBoost model
X, y = df[xvars], df['Big4']
```

```
In [20]: sns.heatmap(df[xvars].corr())
```

```
Out[20]: <Axes: >
```



```
In [23]: pip install scikit-learn
```

```
Collecting scikit-learn
  Downloading scikit_learn-1.3.2-cp310-cp310-macosx_10_9_x86_64.whl (10.2 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 10.2/10.2 MB 21.0 MB/s eta 0:00:
0000:010:01
Collecting scipy>=1.5.0
  Downloading scipy-1.11.3-cp310-cp310-macosx_10_9_x86_64.whl (37.3 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 37.3/37.3 MB 29.2 MB/s eta 0:00:
0000:0100:01
Requirement already satisfied: joblib>=1.1.1 in ./anaconda3/lib/python3.10/site-packages (from scikit-learn) (1.1.1)
Requirement already satisfied: numpy<2.0,>=1.17.3 in ./anaconda3/lib/python3.10/site-packages (from scikit-learn) (1.23.5)
Requirement already satisfied: threadpoolctl>=2.0.0 in ./anaconda3/lib/python3.10/site-packages (from scikit-learn) (2.2.0)
Installing collected packages: scipy, scikit-learn
Successfully installed scikit-learn-1.3.2 scipy-1.11.3
Note: you may need to restart the kernel to use updated packages.
```

In [29]: `pip install xgboost`

```
Collecting xgboost
  Downloading xgboost-2.0.1-py3-none-macosx_10_15_x86_64.macosx_11_0_x86_64.macosx_12_0_x86_64.whl (2.2 MB)
  ━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━━ 2.2/2.2 MB 6.7 MB/s eta 0:00:00a
0:00:01
Requirement already satisfied: scipy in ./anaconda3/lib/python3.10/site-packages (from xgboost) (1.11.3)
Requirement already satisfied: numpy in ./anaconda3/lib/python3.10/site-packages (from xgboost) (1.23.5)
Installing collected packages: xgboost
Successfully installed xgboost-2.0.1
Note: you may need to restart the kernel to use updated packages.
```

In [30]: `from sklearn.model_selection import train_test_split`
`import xgboost as xgb`

In [31]: `Y = df['Big4']`
`X = df[xvars]`
`X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2, random_state=42)`
`xgb_model = xgb.XGBRegressor(random_state=42)`
`xgb_model.fit(X_train, Y_train)`

Out[31]:


```

XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
```


In [56]: `# The SHAP Values`
`import shap`

```
explainer = shap.Explainer(xgb_model)
shap_values = explainer(X_test)
shap.plots.bar(shap_values, max_display=10) # default is max_display=12
```


[17:24:22] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:1240: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.
 [17:24:23] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:1240: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.




```
In [57]: shap.plots.bar(shap_values.cohorts(2).abs.mean(0)) # Cohort plot
```




```
In [58]: fig = plt.gcf() # gcf means "get current figure"
fig.set_figheight(11)
fig.set_figwidth(11)
shap.plots.heatmap(shap_values[1:100])
```




```
In [59]: shap.plots.heatmap(shap_values[200:300])
```




```
In [60]: shap.plots.waterfall(shap_values[0]) # For the first observation
```


```
In [61]: shap.plots.waterfall(shap_values[1]) # For the second observation
```



```
In [62]: shap.plots.bar(shap_values[0]) # For the first observation barplot
```

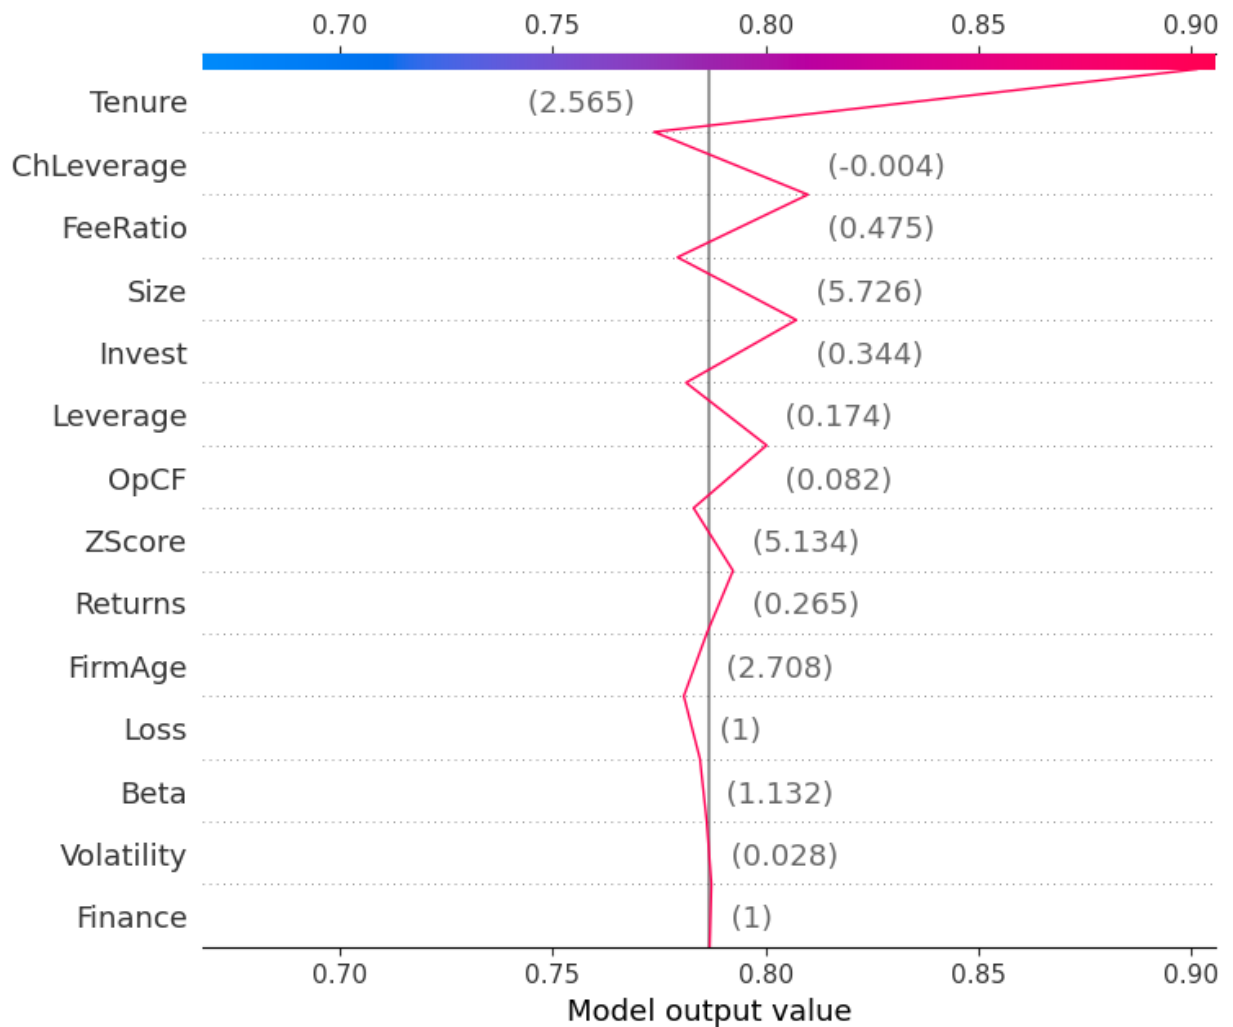


```
In [63]: shap.plots.bar(shap_values[1]) # For the second observation barplot
```



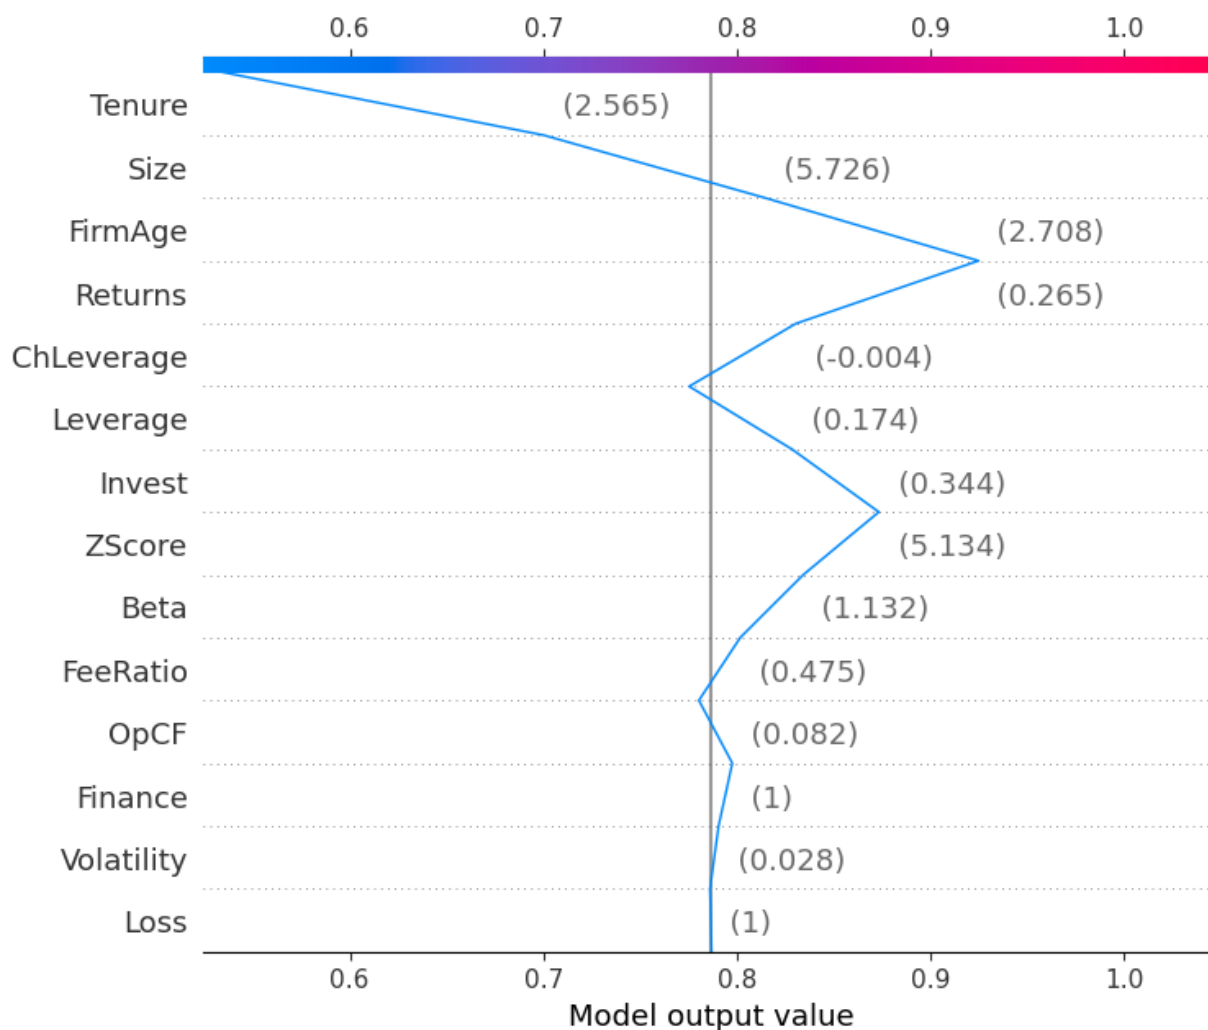
```
In [64]: # Decision plot
expected_value = explainer.expected_value
#print("The expected value is ", expected_value)
print("The final prediction is ", xgb_model.predict(X_test)[0])
shap_values = explainer.shap_values(X_test)[0]
shap.decision_plot(expected_value, shap_values, X_test)
```

The final prediction is 0.9011796



```
In [65]: shap_values = explainer.shap_values(X_test)[1]
# print("The expected value is ", expected_value)
print("The final prediction is ", xgb_model.predict(X_test)[1])
shap.decision_plot(expected_value, shap_values, X_test)
```

The final prediction is 0.5346452



```
In [66]: # Binary Target
xgb_binary_model = xgb.XGBRegressor(objective='reg:logistic', random_state=42)
xgb_binary_model.fit(X_train, Y_train)
Y_pred = xgb_binary_model.predict(X_train)[0]
Y_pred
```

Out[66]: 0.99997056

```
In [67]: # Waterfall plot
explainer = shap.Explainer(xgb_binary_model)
xgb_binary_shap_values = explainer(X_train)
```

[17:27:00] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:124
0: Saving into deprecated binary model format, please consider using `json` or
`ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.
[17:27:01] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:124
0: Saving into deprecated binary model format, please consider using `json` or
`ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.

```
In [68]: from scipy.special import expit
expit(9.676)
```

Out[68]: 0.9999372318494586

```
In [69]: original_shap_values = xgb_binary_shap_values
original_shap_values
```

```

Out[69]: .values =
array([[ 1.5777497 ,  0.34928417, -0.38742083, ...,  0.35079858,
         0.06024942, -0.14367965],
       [ 2.7337563 ,  1.8042798 , -0.5370287 , ...,  0.1392942 ,
         0.13485812,  0.10250279],
       [-1.4035889 ,  0.03870981, -1.2086381 , ..., -0.15479349,
        -0.35088238, -0.23727198],
       ...,
       [-0.5556149 , -0.2396371 , -1.6877143 , ..., -0.16619559,
        -0.14312561,  0.05530791],
       [ 1.8892435 ,  0.5555458 , -0.41296577, ..., -0.2661265 ,
         0.98607045, -0.00834135],
       [-1.0717726 , -0.41847277,  1.1942096 , ..., -0.26563877,
         0.5042884 ,  0.8585152  ]], dtype=float32)

.base_values =
array([1.5671108, 1.5671108, 1.5671108, ..., 1.5671108, 1.5671108,
       1.5671108], dtype=float32)

.data =
array([[2.3025851 , 0.05320805, 2.99573231, ..., 0.27531603, 0.96938528,
        0.01808455],
       [2.8903718 , 0.36467293, 3.17805386, ..., 0.45569211, 1.44153595,
        0.01213791],
       [0.6931472 , 0.22529149, 3.13549423, ..., 0.40078792, 0.78609472,
        0.01260973],
       ...,
       [1.9459101 , 0.05312637, 2.94443893, ..., 0.41147998, 1.55957085,
        0.02640458],
       [2.5649493 , 0.07708553, 3.33220458, ..., 0.72826647, 2.23370677,
        0.03271395],
       [1.0986123 , 0.02189781, 1.38629436, ..., 0.7057573 , 2.23278385,
        0.02153001]])

```

```

In [70]: # Compute the transformed base value, which consists in applying the logit function
from scipy.special import expit # Importing the logit function for the base value
untransformed_base_value = original_shap_values.base_values[-1]
untransformed_base_value

```

```

Out[70]: 1.5671108

```

```

In [71]: original_shap_values.values

```

```

Out[71]: array([[ 1.5777497 ,  0.34928417, -0.38742083, ...,  0.35079858,
         0.06024942, -0.14367965],
       [ 2.7337563 ,  1.8042798 , -0.5370287 , ...,  0.1392942 ,
         0.13485812,  0.10250279],
       [-1.4035889 ,  0.03870981, -1.2086381 , ..., -0.15479349,
        -0.35088238, -0.23727198],
       ...,
       [-0.5556149 , -0.2396371 , -1.6877143 , ..., -0.16619559,
        -0.14312561,  0.05530791],
       [ 1.8892435 ,  0.5555458 , -0.41296577, ..., -0.2661265 ,
         0.98607045, -0.00834135],
       [-1.0717726 , -0.41847277,  1.1942096 , ..., -0.26563877,
         0.5042884 ,  0.8585152  ]], dtype=float32)

```

```

In [76]: import numpy as np
original_explanation_distance = np.sum(original_shap_values.values, axis=1)[2]

```

```
original_explanation_distance
```

```
Out[76]: -4.2045155
```

```
In [78]: distance_to_explain = abs(Y_pred - untransformed_base_value)
distance_to_explain
```

```
Out[78]: 0.5671402
```

```
In [80]: distance_coefficient = np.abs(original_explanation_distance / distance_to_explain)
distance_coefficient
```

```
Out[80]: 7.413538
```

```
In [81]: original_shap_values.values[2]
```

```
Out[81]: array([-1.4035889e+00,  3.8709812e-02, -1.2086381e+00, -1.5047621e+00,
        1.7621313e+00, -4.8217103e-02, -6.5158802e-01,  1.2490553e-03,
        -2.8659341e-01, -2.4763167e-02, -1.3550682e-01, -1.5479349e-01,
        -3.5088238e-01, -2.3727198e-01], dtype=float32)
```


```
In [82]: shap_values_transformed = original_shap_values / distance_coefficient
shap_values_transformed
```

```
Out[82]: .values =
array([[ 0.21282008,  0.04711437, -0.05225856, ...,  0.04731864,
        0.00812695, -0.01938071],
       [ 0.3687519 ,  0.24337634, -0.07243892, ...,  0.01878917,
        0.01819079,  0.01382643],
       [-0.1893278 ,  0.0052215 , -0.1630312 , ..., -0.02087984,
        -0.04732995, -0.03200523],
       ...,
       [-0.07494599, -0.03232426, -0.22765303, ..., -0.02241785,
        -0.01930598,  0.00746039],
       [ 0.25483695,  0.07493667, -0.05570428, ..., -0.03589737,
        0.13300943, -0.00112515],
       [-0.14456965, -0.05644711,  0.16108498, ..., -0.03583158,
        0.06802263,  0.11580371]], dtype=float32)

.base_values =
array([0.21138501, 0.21138501, 0.21138501, ..., 0.21138501, 0.21138501,
       0.21138501], dtype=float32)

.data =
array([[0.31059193, 0.00717715, 0.40408943, ..., 0.03713693, 0.13075879,
        0.0024394 ],
       [0.38987752, 0.04919013, 0.42868248, ..., 0.06146756, 0.19444642,
        0.00163726],
       [0.09349749, 0.0303892 , 0.42294168, ..., 0.05406163, 0.10603503,
        0.00170091],
       ...,
       [0.26248063, 0.00716613, 0.39717055, ..., 0.05550386, 0.21036796,
        0.00356167],
       [0.34598181, 0.01039794, 0.44947562, ..., 0.09823467, 0.30130105,
        0.00441273],
       [0.14819001, 0.00295376, 0.18699498, ..., 0.09519845, 0.30117656,
        0.00290415]])
```

```
In [83]: base_value = shap_values_transformed.base_values
shap_values_transformed.data = original_shap_values.data
shap.plots.waterfall(shap_values_transformed[0])
```



```
In [85]: def xgb_shap_transform_scale(original_shap_values, Y_pred, which):
    from scipy.special import expit

    # Compute the transformed base value, which consists in applying the logit
    from scipy.special import expit #Importing the logit function for the base
    untransformed_base_value = original_shap_values.base_values[-1]

    # Computing the original_explanation_distance to construct the distance_coe
    original_explanation_distance = np.sum(original_shap_values.values, axis=1)

    base_value = expit(untransformed_base_value) # = 1 / (1+ np.exp(-untransfo
    # Computing the distance between the model_prediction and the transformed b
    distance_to_explain = Y_pred[which] - base_value
    # The distance_coefficient is the ratio between both distances which will b
    distance_coefficient = original_explanation_distance / distance_to_explain
    # Transforming the original shapley values to the new scale
    shap_values_transformed = original_shap_values / distance_coefficient
    # Finally resetting the base_value as it does not need to be transformed
    shap_values_transformed.base_values = base_value
    shap_values_transformed.data = original_shap_values.data

    # Now returning the transformed array
    return shap_values_transformed
```

```
In [90]: obs = 0
Y_pred = xgb_binary_model.predict(X_train)
print("The prediction is ", Y_pred[obs])
```

```
shap_values_transformed = xgb_shap_transform_scale(xgb_binary_shap_values, Y_pred)
shap.plots.waterfall(shap_values_transformed[obs])
```

The prediction is 0.99997056

```
-----
ValueError                                Traceback (most recent call last)
Cell In[90], line 4
      2 Y_pred = xgb_binary_model.predict(X_train)
      3 print("The prediction is ", Y_pred[obs])
----> 4 shap_values_transformed = xgb_shap_transform_scale(xgb_binary_shap_val
ues, Y_pred, 0)
      5 shap.plots.waterfall(shap_values_transformed[obs])

Cell In[85], line 17, in xgb_shap_transform_scale(original_shap_values, Y_pred, which)
      15 distance_coefficient = original_explanation_distance / distance_to_explain
      16 # Transforming the original shapley values to the new scale
----> 17 shap_values_transformed = original_shap_values / distance_coefficient
      18 # Finally resetting the base_value as it does not need to be transformed
      19 shap_values_transformed.base_values = base_value


File ~/anaconda3/lib/python3.10/site-packages/shap/_explanation.py:495, in Explanation.__truediv__(self, other)
      494 def __truediv__(self, other):
--> 495     return self._apply_binary_operator(other, operator.truediv, "__truediv__")

File ~/anaconda3/lib/python3.10/site-packages/shap/_explanation.py:469, in Explanation._apply_binary_operator(self, other, binary_op, op_name)
      467     new_exp.base_values = binary_op(new_exp.base_values, other.base_values)
      468 else:
--> 469     new_exp.values = binary_op(new_exp.values, other)
      470     if new_exp.data is not None:
      471         new_exp.data = binary_op(new_exp.data, other)


ValueError: operands could not be broadcast together with shapes (4000,14) (4000,)
```

```
In [91]: explainer = shap.Explainer(xgb_model)
shap_values = explainer(X_test)
shap.plots.bar(shap_values.cohorts(3).abs.mean(0))
```

```
[16:21:22] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:1240: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.
```


```
In [92]: shap.plots.bar(shap_values.cohorts(3).abs.mean(0), show=False)
fig = plt.gcf() # gcf means "get current figure"
fig.set_figheight(11)
fig.set_figwidth(9)
#plt.rcParams['font.size'] = '12'
ax = plt.gca() #gca means "get current axes"
leg = ax.legend(bbox_to_anchor=(0., 1.02, 1., .102))
for l in leg.get_texts(): l.set_text(l.get_text().replace('Class', 'Klasse'))
plt.show()
```



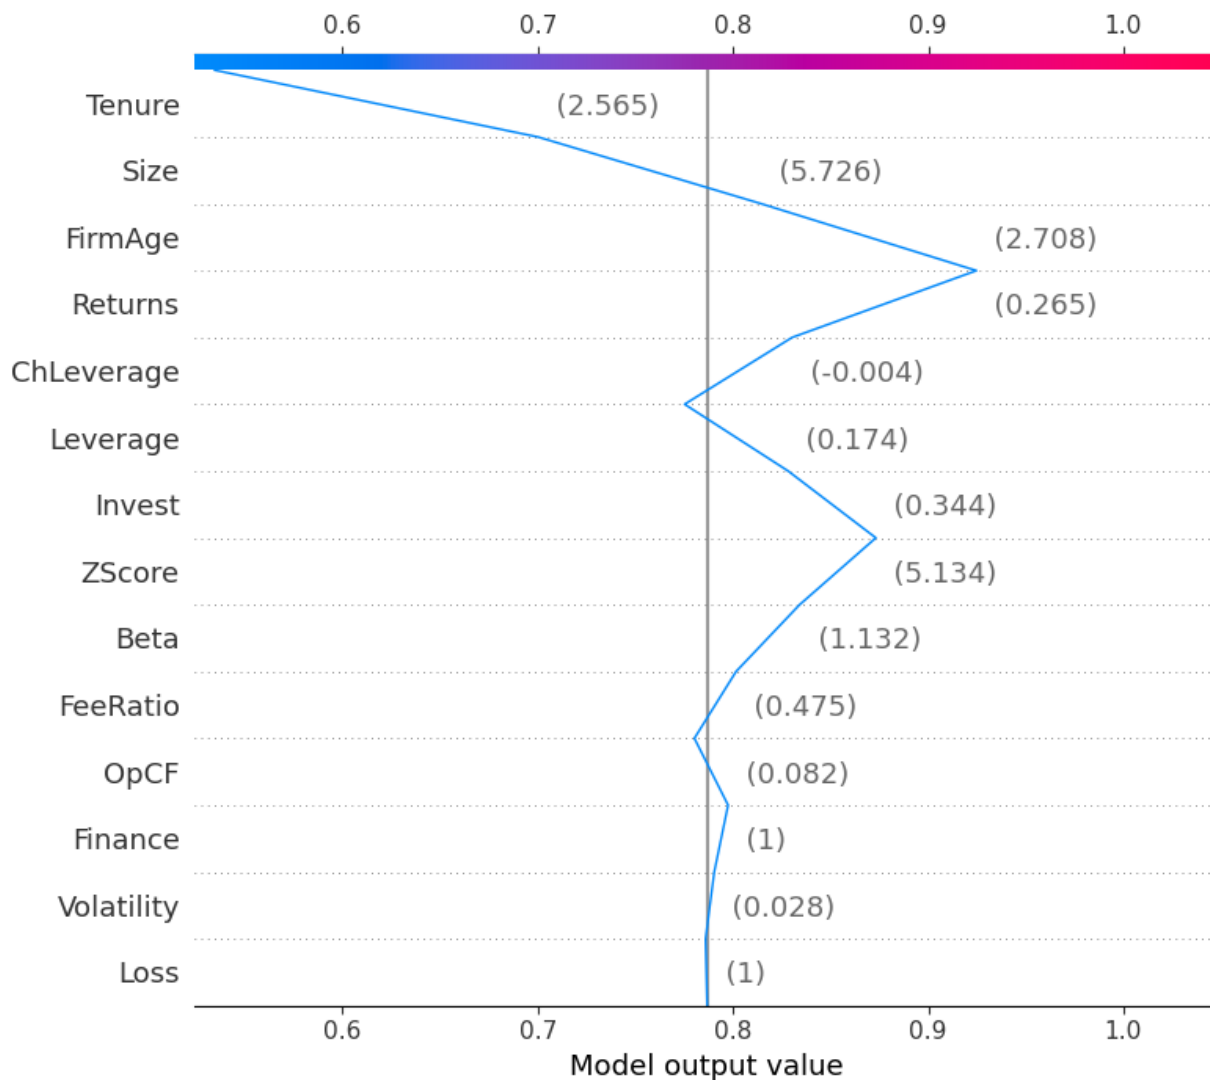
```
In [93]: fig = plt.figure(figsize=(10,5))
```

<Figure size 1000x500 with 0 Axes>

```
In [94]: ax1 = fig.add_subplot(121)
shap_values = explainer.shap_values(X_test)[0]
shap.decision_plot(expected_value, shap_values, X_test, show=False)
ax1.title.set_text('The First Observation')
```




```
In [95]: ax2 = fig.add_subplot(122)
shap_values = explainer.shap_values(X_test)[1]
shap.decision_plot(expected_value, shap_values, X_test, show=False)
ax2.title.set_text('The Second Observation')
plt.tight_layout()
plt.show()
```




```
In [96]: # train an XGBoost model
X, y = df[xvars], df['Big4']
model = xgb.XGBRegressor().fit(X, y)
# explain the model's predictions using SHAP
# (same syntax works for LightGBM, CatBoost, scikit-learn, transformers, Spark,
explainer = shap.Explainer(model)
shap_values = explainer(X)
```

[16:23:13] WARNING: /Users/runner/work/xgboost/xgboost/src/c_api/c_api.cc:124
0: Saving into deprecated binary model format, please consider using `json` or `ubj`. Model format will default to JSON in XGBoost 2.2 if not specified.

```
In [97]: # visualize the first prediction's explanation
shap.plots.waterfall(shap_values[0])
```

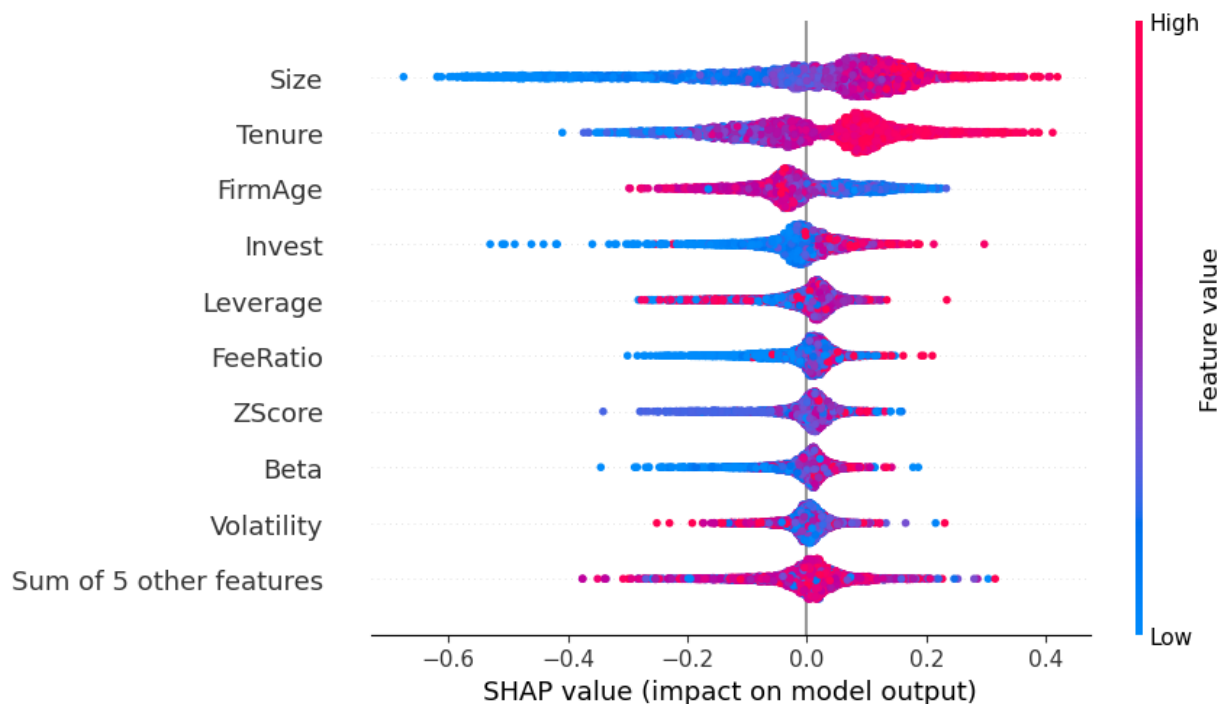


```
In [98]: shap.plots.scatter(shap_values[:, "Size"], color=shap_values)
```

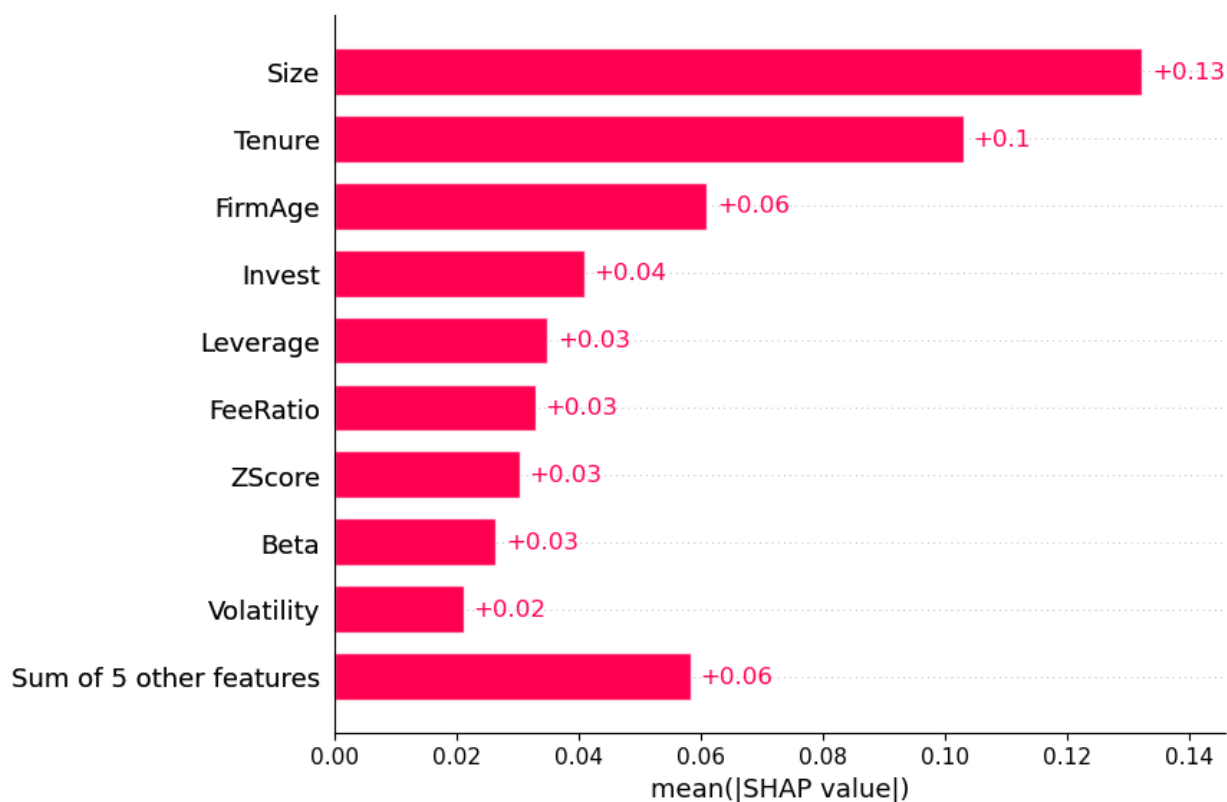


```
In [99]: shap.plots.beeswarm(shap_values)
```

No data for colormapping provided via 'c'. Parameters 'vmin', 'vmax' will be ignored




```
In [100... shap.plots.bar(shap_values)
```



```
In [101... shap.summary_plot(shap_values, df[xvars])
```

No data for colormapping provided via 'c'. Parameters 'vmin', 'vmax' will be ignored



```
In [102... print(df.describe().transpose())  
df.describe().T
```

	count	mean	std	min	25%	50%	
75%	max						
GCO	5000.0	0.010600	0.102419	0.000000e+00	0.000000	0.000000	0.00
0000	1.000000						
Tenure	5000.0	1.904949	0.751237	0.000000e+00	1.386294	2.079442	2.48
4907	2.890372						
Big4	5000.0	0.778600	0.415231	0.000000e+00	1.000000	1.000000	1.00
0000	1.000000						
FeeRatio	5000.0	0.141429	0.138196	0.000000e+00	0.024196	0.104314	0.22
2111	0.621101						
FirmAge	5000.0	2.730144	0.752482	0.000000e+00	2.197225	2.772589	3.21
8876	4.174387						
ZScore	5000.0	2.232858	6.745728	-1.471205e+02	0.484653	1.544287	3.15
1230	100.537613						
Size	5000.0	6.898541	2.108985	1.058137e+00	5.368144	6.890536	8.31
6711	11.736827						
Loss	5000.0	0.483600	0.499781	0.000000e+00	0.000000	0.000000	1.00
0000	1.000000						
Leverage	5000.0	0.584121	0.283002	2.178798e-02	0.384587	0.590995	0.76
2653	2.729931						
ChLeverage	5000.0	0.016008	0.157074	-6.009547e+00	-0.023092	0.007548	0.04
9285	2.395992						
OpCF	5000.0	0.017839	0.189054	-3.164750e+00	-0.002316	0.052809	0.09
7786	0.359912						
Finance	5000.0	0.925400	0.262771	0.000000e+00	1.000000	1.000000	1.00
0000	1.000000						
Invest	5000.0	0.320069	0.299344	1.940000e-07	0.074824	0.209023	0.51
2577	0.995753						
Returns	5000.0	0.414332	0.195552	3.283552e-02	0.296476	0.394603	0.48
8119	1.461594						
Beta	5000.0	1.136815	0.571122	-4.444910e-01	0.750049	1.108556	1.47
3767	2.815003						
Volatility	5000.0	0.025848	0.015291	7.749172e-03	0.015060	0.022943	0.03
2470	0.137875						

Out[102]:

	count	mean	std	min	25%	50%	75%	
GCO	5000.0	0.010600	0.102419	0.000000e+00	0.000000	0.000000	0.000000	1.0
Tenure	5000.0	1.904949	0.751237	0.000000e+00	1.386294	2.079442	2.484907	2.8
Big4	5000.0	0.778600	0.415231	0.000000e+00	1.000000	1.000000	1.000000	1.0
FeeRatio	5000.0	0.141429	0.138196	0.000000e+00	0.024196	0.104314	0.222111	0.
FirmAge	5000.0	2.730144	0.752482	0.000000e+00	2.197225	2.772589	3.218876	4.
ZScore	5000.0	2.232858	6.745728	-1.471205e+02	0.484653	1.544287	3.151230	100.5
Size	5000.0	6.898541	2.108985	1.058137e+00	5.368144	6.890536	8.316711	11.7
Loss	5000.0	0.483600	0.499781	0.000000e+00	0.000000	0.000000	1.000000	1.0
Leverage	5000.0	0.584121	0.283002	2.178798e-02	0.384587	0.590995	0.762653	2.7
ChLeverage	5000.0	0.016008	0.157074	-6.009547e+00	-0.023092	0.007548	0.049285	2.3
OpCF	5000.0	0.017839	0.189054	-3.164750e+00	-0.002316	0.052809	0.097786	0.3
Finance	5000.0	0.925400	0.262771	0.000000e+00	1.000000	1.000000	1.000000	1.0
Invest	5000.0	0.320069	0.299344	1.940000e-07	0.074824	0.209023	0.512577	0.9
Returns	5000.0	0.414332	0.195552	3.283552e-02	0.296476	0.394603	0.488119	1.4
Beta	5000.0	1.136815	0.571122	-4.444910e-01	0.750049	1.108556	1.473767	2.8
Volatility	5000.0	0.025848	0.015291	7.749172e-03	0.015060	0.022943	0.032470	0.7

```
In [105... X, y = df[xvars],df['GCO']
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.25, ran
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.fit_transform(X_test)
from sklearn.linear_model import LogisticRegression
classifier = LogisticRegression(random_state=0)
classifier.fit(X_train, y_train)
# print(classifier.predict(sc.transform([[30,87000]])))
y_pred=classifier.predict(X_test)
print(np.concatenate((y_pred,y_test)))

[0 0 0 ... 0 0 0]
```

```
In [106... model = LogisticRegression(solver='liblinear', C=10, random_state=0).fit(X, y)
model.fit(X, y)
```

Out[106]:

▼

LogisticRegression
LogisticRegression(C=10, random_state=0, solver='liblinear')

```
In [107... model.intercept_
```

Out[107]: array([-0.06138551])

```
In [108... model.coef_
```

```
Out[108]: array([[ 1.15792902e-01, -3.93672366e-01, -6.56029395e-01,
                -4.24949993e-03, -4.55376186e-01,  1.64317434e+00,
                7.42275404e-01,  4.60870433e-02, -1.59900583e+00,
                -4.63880680e-01, -9.23187666e-01, -4.41882120e+00,
                -4.47465708e-01,  3.67202122e+00]])
```

```
In [109... model.predict_proba(X)
```

```
Out[109]: array([[9.99874050e-01, 1.25950001e-04],
                [9.99424928e-01, 5.75071974e-04],
                [9.91115991e-01, 8.88400890e-03],
                ...,
                [9.98976165e-01, 1.02383506e-03],
                [9.99815810e-01, 1.84190342e-04],
                [9.99889245e-01, 1.10755283e-04]])
```

```
In [110... model.score(X, y)
```

```
Out[110]: 0.9896
```

```
In [111... import statsmodels.api as sm
model = sm.Logit(y, X)
result = model.fit(method='newton')
result.summary()
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
Cell In[111], line 1
----> 1 import statsmodels.api as sm
      2 model = sm.Logit(y, X)
      3 result = model.fit(method='newton')

ModuleNotFoundError: No module named 'statsmodels'
```

```
In [112... X = sm.add_constant(X) # add a constant / intercept
model = sm.Logit(y, X)
result = model.fit(method='newton')
result.summary()
```

```
-----
NameError                                          Traceback (most recent call last)
Cell In[112], line 1
----> 1 X = sm.add_constant(X) # add a constant / intercept
      2 model = sm.Logit(y, X)
      3 result = model.fit(method='newton')

NameError: name 'sm' is not defined
```

```
In [114... from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from sklearn.metrics import confusion_matrix, classification_report
params = {'criterion':['gini','entropy'],
          'max_depth':[10,50,100,None],
          'random_state':[123],
          'class_weight':[None,'balanced']}
dt = DecisionTreeClassifier() # Leave this as default
train,test = train_test_split(df,train_size=0.80, random_state=123)
# parameters: n_jobs = -1 - parallelize on all cores; cv=5 - 5-fold cross-val;
# n_iter = 10 - 10 different parameter sets; verbose = 5 - print
# some output
```

```
rs = RandomizedSearchCV(dt,param_distributions=params,n_jobs=-1,cv=5,n_iter=10  
res = rs.fit(train[xvars],train['Big4'])
```

Fitting 5 folds for each of 10 candidates, totalling 50 fits

```
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:  
https://scikit-learn.org/stable/model\_persistence.html#security-maintainabilit  
y-limitations  
warnings.warn(  
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:  
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier  
from version 1.3.2 when using version 1.3.0. This might lead to breaking code  
or invalid results. Use at your own risk. For more info please refer to:
```

```
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
```

```
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
```

```
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347: InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations)
```



```
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to: https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
```



```
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
warnings.warn(
/Users/weiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifie
r from version 1.3.2 when using version 1.3.0. This might lead to breaking cod
e or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainabilit
y-limitations
```

[illegible]


```

r from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(
/Users/yeiserross/anaconda3/lib/python3.10/site-packages/sklearn/base.py:347:
InconsistentVersionWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.3.2 when using version 1.3.0. This might lead to breaking code or invalid results. Use at your own risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations
warnings.warn(

```

```

In [ ]: print("Confusion matrix (Rows = True 0/1, Columns = Predicted 0/1):")
print(confusion_matrix(test['Big4'],rs.predict(test[xvars])))
print("\n\nscikit-learn's classification report, which gives information on accuracy")
print(classification_report(test['Big4'],rs.predict(test[xvars])))
print("\n\nThe search identified this as the best parameters:")
print(res.best_params_)

```

```

In [ ]:

```